

# Data Management Using Tidyverse

Myung Jung Kim

8/7/2023

# Introduction

- **Tidyverse** is a powerful data management tool. It's a collection of R packages like `ggplot2`, `dplyr`, `tidyr`, `stringr` – which include tools for data manipulation, visualization and string manipulation. In this section, we focus on data management tools (mainly `dplyr` and `tidyr`).
- Using tidyverse, we will learn how to do:
  - ① Selecting and Sorting Columns
  - ② Data Reduction and Expansion
  - ③ Grouping and Summarizing
- Basic grammar: **Pipe operator (`%>%`)** combines various functions without the result to a new object.
  - ▶ `data %>% verb 1 %>% verb 2 %>% verb 3`

## Open Demo Data sets: *gapminder*, *diamond*

```
## Gapminder dataset (in gapminder package)
install.packages("gapminder")
library(gapminder)
data(gapminder) # load data
names(gapminder) # check the variable names
View(gapminder) # open a data viewer
```

```
## Diamond dataset (in tidyverse package)
install.packages("tidyverse")
library(tidyverse)
data(diamonds) # load data
names(diamonds) # check the variable names
View(diamonds) # open a data viewer
```

# Section 1

## Selecting and Arranging Columns

## 1-1. select() function

- Selecting certain columns

- ▶ gapminder has 6 variables: country, continent, year, lifeExp, pop, gdpPercap.
- ▶ Let's say we only need 3 columns: country, year, gdpPercap.

```
## Put data name, %>% (pipe), and select() function
```

```
gapminder %>%  
  select(country, year, gdpPercap)
```

```
## You can save it as a separate dataset (gapminder2)
```

```
gapminder2 <- gapminder %>%  
  select(country, year, gdpPercap)
```

- De-selecting certain columns
  - ▶ You **de-select** certain columns using `select()` with **minus sign (-)**.
  - ▶ let's say you don't need 'continent' and 'pop' from the `gapminder`.

```
## De-select continent, pop
```

```
gapminder %>%  
  select(-continent, -pop)
```

```
## Save it as a separate dataset (gapminder3)
```

```
gapminder3 <- gapminder %>%  
  select(-continent, -pop)
```

- `everything()` function is useful when you want to locate few specific variables to the front and leave everything else as now.

```
## move lifeExp to the front
```

```
gapminder %>%  
  select(lifeExp, everything())
```

## Exercise 1-1. Selecting Columns

- ① How many variables do we have in the diamonds dataset?
- ② Can you list the names?
- ③ **Select** columns 'carat', 'color', 'clarity' from diamonds and save it as a separate dataset called diamond2.
- ④ **De-select** columns 'x', 'y', 'z' from the diamonds dataset.
- ⑤ Move 'x', 'y', 'z' to the very front from the original diamonds dataset and leave every other columns as the order of now.

## 1-2. arrange() function

- `arrange()` helps reorder the rows of data by using column names.
- Right now, `gapminder`'s column orders by “**country**”, alphabetically.

country <fctr>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>	gdpPercap <dbl>
Afghanistan	Asia	1952	28.801	8425333	779.4453
Afghanistan	Asia	1957	30.332	9240934	820.8530
Afghanistan	Asia	1962	31.997	10267083	853.1007
Afghanistan	Asia	1967	34.020	11537966	836.1971
Afghanistan	Asia	1972	36.088	13079460	739.9811

- If you want to rearrange the data by “**year**”:

```
gapminder %>% arrange(year)
```

country <fctr>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>	gdpPercap <dbl>
Afghanistan	Asia	1952	28.801	8425333	779.4453
Albania	Europe	1952	55.230	1282697	1601.0561
Algeria	Africa	1952	43.077	9279525	2449.0082
Angola	Africa	1952	30.015	4232095	3520.6103
Argentina	Americas	1952	62.485	17876956	5911.3151



- In R, the default setting arranges data in **ascending** order (either alphabet or number). If you want to arrange data in **descending** order, you need to specify it using `desc()`.

```
gapminder %>% arrange(desc(country))
```

country <fctr>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>
Zimbabwe	Africa	1952	48.451	3080907
Zimbabwe	Africa	1957	50.469	3646340
Zimbabwe	Africa	1962	52.358	4277736
Zimbabwe	Africa	1967	53.995	4995432
Zimbabwe	Africa	1972	55.635	5861135

```
gapminder %>% arrange(desc(year))
```

country <fctr>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>
Afghanistan	Asia	2007	43.828	31889923
Albania	Europe	2007	76.423	3600523
Algeria	Africa	2007	72.301	33333216
Angola	Africa	2007	42.731	12420476
Argentina	Americas	2007	75.320	40301927

- `rename()` function: it's to change the name of a column.
  - ▶ `data %>% rename(NEW name = OLD column name)`

```
gapminder %>%  
  rename (state = country)
```

## Exercise 1-2. Sorting Columns

- ① Arrange diamonds dataset by a column **cut**.
- ② Arrange diamonds dataset by **carat** in descending order.
- ③ Rename the column 'clarity' into 'Clarity'.

## Section 2

### Data Reduction and Expansion

## 2-1. filter() function: conditional filtering

- You use filter when you want to look only at a subset of your observations, based on a particular condition.

```
gapminder %>%  
  filter(year == 2002)
```

A tibble: 142 × 6

country <fctr>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>	gdpPercap <dbl>
Afghanistan	Asia	2002	42.129	25268405	726.7341
Albania	Europe	2002	75.651	3508512	4604.2117
Algeria	Africa	2002	70.994	31287142	5288.0404
Angola	Africa	2002	41.003	10866106	2773.2873
Argentina	Americas	2002	74.340	38331121	8797.6407
Australia	Oceania	2002	80.370	19546792	30687.7547
Austria	Europe	2002	78.980	8148312	32417.6077
Bahrain	Asia	2002	74.795	656397	23403.5593
Bangladesh	Asia	2002	62.013	135656790	1136.3904
Belgium	Europe	2002	78.320	10311970	30485.8838

- Filtering condition **“A and B”**: we can specify multiple conditions in the filter. Each of the conditions is separated by a comma:

```
gapminder %>%  
  filter(year == 2002, continent == "Africa")
```

A tibble: 52 × 6

country <fctr>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>	gdpPercap <dbl>
Algeria	Africa	2002	70.994	31287142	5288.0404
Angola	Africa	2002	41.003	10866106	2773.2873
Benin	Africa	2002	54.406	7026113	1372.8779
Botswana	Africa	2002	46.634	1630347	11003.6051
Burkina Faso	Africa	2002	50.650	12251209	1037.6452
Burundi	Africa	2002	47.360	7021078	446.4035
Cameroon	Africa	2002	49.856	15929988	1934.0114
Central African Republic	Africa	2002	43.308	4048013	738.6906
Chad	Africa	2002	50.525	8835739	1156.1819
Comoros	Africa	2002	62.974	614382	1075.8116

- Filtering condition “**A or B**”: we can specify multiple conditions in the filter. Each of the conditions is separated by a “|” :

```
gapminder %>%  
  filter(year == 2002 | continent == "Africa")
```

A tibble: 714 × 6

country <fctr>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>	gdpPercap <dbl>
Afghanistan	Asia	2002	42.129	25268405	726.7341
Albania	Europe	2002	75.651	3508512	4604.2117
Algeria	Africa	1952	43.077	9279525	2449.0082
Algeria	Africa	1957	45.685	10270856	3013.9760
Algeria	Africa	1962	48.303	11000948	2550.8169
Algeria	Africa	1967	51.407	12760499	3246.9918
Algeria	Africa	1972	54.518	14760787	4182.6638
Algeria	Africa	1977	58.014	17152804	4910.4168
Algeria	Africa	1982	61.368	20033753	5745.1602
Algeria	Africa	1987	65.799	23254956	5681.3585

- You can summarize information of the filtered data.

*## What's the mean life expectancy in Africa in year 2002?*

```
gapminder %>%  
  filter(year == 2002, continent == "Africa")%>%  
  summarize(mean(lifeExp))
```

```
mean(lifeExp)  
      <dbl>
```

---

```
53.32523
```

---



## 2-2. slice() function:

- Filtering rows works well when we know the data columns and values by name. But if that information is not as readily available, we can reduce the data using slice() functions:
  - ▶ `%>%slice(100:300)`: slice rows 100 to 300
  - ▶ `%>%slice_head(n= 10)`: show the first n rows
  - ▶ `%>%slice_tail(n= 10)`: show the last n rows
  - ▶ `%>%slice_min(column name)`: show the smallest value of the column
  - ▶ `%>%slice_max(column name)`: show the largest value of the column

```
gapminder %>%
  slice_head(n=8)
```

A tibble: 8 × 6

country <fctr>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>	gdpPercap <dbl>
Afghanistan	Asia	1952	28.801	8425333	779.4453
Afghanistan	Asia	1957	30.332	9240934	820.8530
Afghanistan	Asia	1962	31.997	10267083	853.1007
Afghanistan	Asia	1967	34.020	11537966	836.1971
Afghanistan	Asia	1972	36.088	13079460	739.9811
Afghanistan	Asia	1977	38.438	14880372	786.1134
Afghanistan	Asia	1982	39.854	12881816	978.0114
Afghanistan	Asia	1987	40.822	13867957	852.3959

```
gapminder %>%
  slice_max(pop)
```

A tibble: 1 × 6

country <fctr>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>	gdpPercap <dbl>
China	Asia	2007	72.961	1318683096	4959.115

## 2-3. dropping NAs: drop\_na()

- If you don't know a value is missing or not, you can use the base R function `is.na()`.
- `drop_na()` function removes NA values.

```
table(is.na(gapminder)) # no NAs.  
# Create a vector with NAs  
v <- data.frame("age"=c(1.2, 4.5, NA, 8.9, NA),  
                "Name" = c("sua", "yul", "alex", "peter", NA))  
# Drop NAs  
v %>% drop_na
```

	age	Name	Description: df [3 × 2]	
1	1.2	sua	age	Name
2	4.5	yul	<dbl>	<chr>
3	NA	alex	1.2	sua
4	8.9	peter	4.5	yul
5	NA	NA	8.9	peter

## 2-4. mutate() function: create columns of new information

- **mutate()** creates new variables (and preserves existing ones).
- **transmute()** adds new variables and drops existing ones.
- Let's make a new dummy variable called **post2000** if **year** is 2000 or later.
- **ifelse(test\_expression, x, y)**: The output vector has the element **x** if the output of the **test\_expression** is **TRUE**. If the output is **FALSE**, then the element in the output vector will be **y**.

```
# New column "post2000" using mutate()
```

```
gapminder %>%
```

```
  mutate(post2000 = ifelse(year >= 2000, 1, 0))
```

A tibble: 1,704 × 7

country <fctr>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>	gdpPercap <dbl>	post2000 <dbl>
Afghanistan	Asia	1952	28.80100	8425333	779.4453	0
Afghanistan	Asia	1957	30.33200	9240934	820.8530	0
Afghanistan	Asia	1962	31.99700	10267083	853.1007	0
Afghanistan	Asia	1967	34.02000	11537966	836.1971	0
Afghanistan	Asia	1972	36.08800	13079460	739.9811	0
Afghanistan	Asia	1977	38.43800	14880372	786.1134	0

## Exercise 2. Data Reduction and Expansion

- 1 From the diamonds dataset, filter rows so that you only observe **Premium cut AND carat size above 0.70**. What's the *mean price* of the diamonds that meet the conditions?
- 2 Print the last 10 rows of the diamond dataset.
- 3 Is there NA in the diamond dataset? If so, drop all NAs and save it as diamonds3.
- 4 Create a new column called "standard" in the diamonds dataset and fill it with "high" if the cut is either Premium or Ideal, and otherwise "low".

## Section 3

### Grouping and Summarizing

## group\_by() function with summarize() function

- If we want to assess information by group, we use **group\_by()** together with **summarize()**.

*## There are 7 color groups in the diamonds dataset.*

*# What's the average price of each color group?*

```
diamonds %>% group_by(color)%>% summarize(mean(price))
```

A tibble: 7 × 2

color <ord>	mean(price) <dbl>
D	3169.954
E	3076.752
F	3724.886
G	3999.136
H	4486.669
I	5091.875
J	5323.818

## group\_by(), mutate(), and ungroup() together

- We could also utilize mutate() after group\_by() to add a new column based on the group.

```
## A new column 'm' for mean price of same color
diamonds %>%
  group_by(color)%>%
  mutate(m = mean(price))%>%
  ungroup() #ungroup after creating new object based on group
```

carat <dbl>	cut <ord>	color <ord>	clarity <ord>	depth <dbl>	table <dbl>	price <int>	x <dbl>	y <dbl>	z <dbl>	m <dbl>
0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43	3076.752
0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31	3076.752
0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31	3076.752
0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63	5091.875
0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75	5323.818
0.24	Very Good	J	VVS2	62.8	57.0	336	3.94	3.96	2.48	5323.818
0.24	Very Good	I	VVS1	62.3	57.0	336	3.95	3.98	2.47	5091.875
0.26	Very Good	H	SI1	61.9	55.0	337	4.07	4.11	2.53	4486.669
0.22	Fair	E	VS2	65.1	61.0	337	3.87	3.78	2.49	3076.752
0.23	Very Good	H	VS1	59.4	61.0	338	4.00	4.05	2.39	4486.669
0.30	Good	J	SI1	64.0	55.0	339	4.25	4.28	2.73	5323.818



## count(), tally()

- When we want to count observations by group, we use **count()** or **tally()**. They are convenient ways to get a sense of the distribution of values in a dataset.

```
diamonds %>% count(color)
diamonds %>% group_by(color)%>% tally() #Same results
```

A tibble: 7 × 2

color <ord>	n <int>
D	6775
E	9797
F	9542
G	11292
H	8304
I	5422
J	2808

```
# You can count based on multiple grouping conditions  
## 7 colors and 5 cuts => 35 groups  
diamonds %>% count(color, cut)
```

A tibble: 35 × 3

<b>color</b> <ord>	<b>cut</b> <ord>	<b>n</b> <int>
D	Fair	163
D	Good	662
D	Very Good	1513
D	Premium	1603
D	Ideal	2834
E	Fair	224
E	Good	933
E	Very Good	2400
E	Premium	2337
E	Ideal	3903

## Exercise 3. Grouping and Summarizing

- 1 From the `gapminder` dataset, what's the mean life expectancy of each continent?
- 2 From the `gapminder` dataset, what's the mean life expectancy of each continent in **year 2007**?
- 3 From the `gapminder` dataset, create a new column called `continent_avg_life` which shows each **continent's average life expectancy** (lifeExp) in the **particular year**. Save it as a new data called `gapminder4`.
- 4 Using `group_by()`, count how many countries belong to each continent. (Hint: use `unique(country)` somewhere)

## ***Any Question?***

*Feel free to reach out for anything :)*

*- Myung Jung Kim (mjkim@illinois.edu)*